START HERE

Scam Spotter - MVP Overview

Scam Spotter checks a pasted SMS, email or chat message, finds visible warning signs using transparent rules, shows a 0-100 risk score with the suspicious phrases highlighted, and tells the user what to do next. This pack contains one specification per MVP feature (14 in total). Each can be handed to whoever builds that part.

Scoring criteria and where to invest

Criterion	Marks	What wins it	Main features
Completeness	35	Whole flow works: message in, result out, next steps shown	1, 2, 3, 4, 7, 8, 9
Monetization	30	Credible target user, pricing, sustainability plan	11 (plus 12, 13 as paid tiers)
Innovation	20	Visible rules and highlighted evidence, not a black box	2, 5, 6, 12
UI / UX	15	Polished, readable, mobile-friendly, easy to understand	10, 4, 5, 6

> One recommendation: your table marks Monetization as P1, but it is worth 30 marks - more than UI/UX and Innovation combined would suggest by effort. It needs mostly thinking and writing, not code. Plan to finish it before the demo.

Feature index

#	Feature	Priority	Effort (rough)
01	Message Input	P0	Small (~1-2 hrs)
02	Rule-Based Detection	P0	Medium (~3-5 hrs)
03	Risk Scoring (0-100)	P0	Small (~1-2 hrs)
04	Risk Classification	P0	Small (~1 hr)
05	Message Highlighting	P0	Medium (~2-4 hrs)
06	Explanation Panel	P0	Medium (~2-3 hrs)
07	Safe Next Steps	P0	Small (~1-2 hrs)
08	Example Scam Message	P0	Small (~30 min)
09	False-Positive Handling	P0	Medium (~3 hrs)
10	Mobile-Responsive Design	P0	Medium (~3-4 hrs, mostly polish)
11	Monetization Plan	P1 (recommend treating as P0)	Medium (~3-4 hrs, mostly thinking and writing)
12	AI-Assisted Analysis	P1	Medium-Large (~4-6 hrs)
13	Scan History	P1	Small-Medium (~2-3 hrs)
14	URL Reputation Checks	P1	Medium (~3-5 hrs)

Suggested build order

- **Step 1 - Skeleton:** message input (1) and example message (8): get 'paste, click, see something' working first.
- **Step 2 - Brain:** rule detection (2), scoring (3), classification (4). Write unit tests for the five example cases as you go.
- **Step 3 - Evidence:** highlighting (5) and explanation panel (6). This is the point where the demo starts to impress.
- **Step 4 - Action:** safe next steps (7) and false-positive handling (9) with the 40-message test set.
- **Step 5 - Polish:** mobile layout (10). Test on a phone-sized screen.
- **Step 6 - Business:** monetization plan (11). Write it while you test.
- **Step 7 - If time permits:** scan history (13), URL reputation (14), AI second opinion (12).

Shared scoring rules (used by several features)

Rule	Points	Note
Fake urgency	+20	From your brief
Account threat	+15	From your brief
Suspicious link	+25	From your brief
OTP request	+40	From your brief
Password / PIN / card request	+40	Proposed; shares a group with OTP so it never double-counts

Levels: Low 0-29, Medium 30-59, High 60-100. Score = sum of fired rules (each once, credential group counted once), capped at 100. A single rule can never exceed 40, so High always needs at least two independent warning signs.

Technical suggestion (beginner friendly)

For the prototype, build the rules engine in plain JavaScript that runs in the browser, inside a single web page. The result: no server to set up, instant answers, nothing uploaded, and nothing that can break on the demo day because of Wi-Fi. A small Python (Flask) server can be added later only for AI analysis and URL reputation, where an API key must stay secret.

Assumptions in this pack

- These specs are built from your written brief and the feature-list text. I could not see the original project-brief images, so confirm that the example message wording and the exact marks match them.
- The example SBI message is fictional and illustrative. The Rs prices, effort times and helpline details are proposals or estimates to verify before presenting.
- No statistics are quoted. If you add market figures to the monetization pitch, cite the source.